

Autonomous Data Science Co-Pilot

Enterprise AI Dataset Analysis Report

Generated on 16 July 2026 15:47

Dataset Summary

Metric	Value
Rows	31
Columns	8
Duplicate Rows	1
Memory Usage	2.71 KB MB

AI Insights

As a Senior Data Scientist, I've analyzed the provided dataset information. Here's a breakdown of my findings and recommendations:

1. Dataset Overview

This dataset consists of 31 rows and 8 columns, providing a snapshot of employee information. It includes unique identifiers (`Employee\_ID`, `Name`), demographic details (`Age`, `City`), professional attributes (`Department`, `Experience\_Years`), and key performance/compensation metrics (`Salary`, `Performance\_Score`).

- * **Size**: Very small (31 records, 8 features).
- * **Column Types**: A mix of numeric (`Employee\_ID`, `Age`, `Salary`, `Experience\_Years`, `Performance\_Score`) and categorical (`Name`, `Department`, `City`) data.
- * **Content**: Focuses on employee demographics, experience, department, and crucial metrics like salary and performance.

2. Data Quality Assessment

The dataset presents several significant data quality issues, particularly concerning missing values and duplicates.

- * **Duplicate Rows**: One duplicate row exists (approx. 3.2% of the dataset), which is a notable proportion for such a small dataset.
- * **Missing Values**:
 - * **Excellent (0% missing)**: `Employee\_ID`, `Name`, `Experience\_Years`. These columns are complete.
 - * **Manageable (3-7% missing)**: `Age` (6.45%), `Department` (3.23%), `City` (3.23%). These percentages are relatively low and can be handled without major data loss.

* **Significant (16% missing)**: `Performance\_Score` (16.13%). This is a substantial amount for a key metric and requires careful imputation or handling.

* **Critical (41.94% missing)**: `Salary`. Almost half of the salary data is missing, making this column highly compromised. This is the most severe data quality issue.

* **Outliers**: No outliers were detected in any of the numeric columns (`Employee\_ID`, `Age`, `Salary`, `Experience\_Years`, `Performance\_Score`). This is positive, suggesting the existing data points are within expected ranges for their respective features.

* **Data Types**: Column types appear appropriate for the data they represent.

3. Potential Problems

Based on the data quality assessment, the following problems are anticipated:

1. **Limited Statistical Power and Generalizability**: The extremely small dataset size (31 rows) inherently limits the robustness of any statistical analysis and the ability of machine learning models to generalize well to unseen data.
2. **Unreliable `Salary` Data**: With nearly 42% missing values, the `Salary` column is severely compromised. Using it directly as a feature or a target variable without aggressive, and potentially biased, imputation will be challenging. Dropping rows with missing `Salary` would reduce the dataset to just 18 rows, making any analysis or modeling highly unstable.
3. **Challenges with `Performance\_Score`**: While less severe than `Salary`, 16% missing values in `Performance\_Score` will require thoughtful imputation to maintain data integrity and avoid bias.
4. **Impact of Duplicate Row**: A single duplicate in such a small dataset can skew descriptive statistics and potentially lead to overfitting if not removed before model training.
5. **`Name` Column Utility**: The `Name` column is likely unique for each employee (high cardinality) and serves as an identifier rather than a useful feature for most machine learning tasks.
6. **Representativeness**: Due to the small size and high missing rates in key columns, it's questionable whether the remaining data is truly representative of the overall employee population.

4. Recommended Preprocessing

To prepare this dataset for analysis and potential machine learning tasks, I recommend the following steps:

1. **Handle Duplicate Rows**:

* **Action**: Remove the single duplicate row. For small datasets, duplicates can disproportionately impact results.

* **Rationale**: Ensures each observation is unique, preventing bias in calculations and model training.

2. **Address Missing Values**:

* **`Salary` (Critical)**:

* **Option 1 (Preferred if `Salary` is the ML target)**: If the primary goal is to predict `Salary`, rows with missing `Salary` should ideally be excluded from the *training* set to avoid imputing the target variable. This will drastically reduce the training data to only 18 samples, which is a major constraint.

* **Option 2 (If `Salary` is a feature)**: If `Salary` is used as a feature, consider a more sophisticated imputation method like K-Nearest Neighbors (KNN) imputation, or even multiple imputation, leveraging correlations with `Experience\_Years`, `Department`, and `Performance\_Score`. Simple mean/median imputation is too simplistic given the high missing rate and could introduce significant bias. Alternatively, create a binary flag indicating "missing salary" and impute with a specific value (e.g., 0 or mean).

* **`Performance\_Score` (Significant)**:

* **Action**: Impute using the median or mean. Given the dataset size, the median is generally more robust to potential subtle skewness, even though no outliers were detected. KNN imputation could also be considered.

* **`Age` (Manageable)**:

* **Action**: Impute with the median. `Age` distributions can sometimes be skewed.

* **`Department`, `City` (Manageable)**:

* **Action**: Impute with the mode (most frequent category) or a distinct "Unknown" category.

* **General Note on Imputation**: For such a small dataset, any imputation method carries a risk of introducing bias or noise. It's crucial to acknowledge this limitation.

3. **Feature Selection/Engineering**:

* **`Name`**: Drop the `Name` column as it's an identifier and likely has too high cardinality for a useful feature.

* **`Employee\_ID`**: Drop this column after verifying its role (if it's just an identifier and not a foreign key for another table).

* **Categorical Encoding**: Encode `Department` and `City` into numerical representations. One-Hot Encoding is suitable for nominal categories, preventing the model from assuming an ordinal relationship.

4. **Feature Scaling (Conditional)**:

* **Action**: Apply a scaling method (e.g., StandardScaler or MinMaxScaler) to numerical features like `Age`, `Experience\_Years`, and `Performance\_Score` (and `Salary` if used as a feature).

* **Rationale**: This is important for many machine learning algorithms (e.g., SVMs, K-Nearest Neighbors, linear models, neural networks) that are sensitive to feature scales. Tree-based models are generally less sensitive.

5. Suggested ML Task

Given the available data and its quality, predicting `Performance\_Score` seems like the most practical and valuable initial machine learning task.

****Suggested Task: Employee Performance Score Prediction (Regression)****

* ****Goal****: Predict an employee's `Performance\_Score` based on other available features such as `Age`, `Department`, `Experience\_Years`, and potentially `Salary` (if reliable imputation can be achieved).

* ****Rationale****:

* `Performance\_Score` is a continuous numeric variable, making it suitable for regression.

* While `Performance\_Score` has missing values (16.13%), this is more manageable than `Salary`.

* Understanding the drivers of performance is highly valuable for talent management, training program development, and resource allocation.

* ****Challenges****: The small dataset size will still be a major limitation for building a robust predictive model. Cross-validation strategies will need to be carefully chosen (e.g., Leave-One-Out Cross-Validation, if feasible, or very small K-fold).

* ****Alternative Task (More challenging due to data quality)****: Employee Salary Prediction (Regression). This would be highly valuable but extremely difficult due to the severe missing values in `Salary`.

6. Business Insights

Despite the data quality challenges, even preliminary analysis or a successful (albeit limited) ML model could provide valuable business insights:

1. ****Identifying Performance Drivers****: Even with a small dataset, initial correlations might emerge. For example, does `Experience\_Years` or a particular `Department` significantly influence `Performance\_Score`? This can inform training programs, career pathing, and hiring profiles.
2. ****Compensation Fairness & Strategy (Conditional on Salary Data)****: If `Salary` data can be made sufficiently reliable, one could analyze if salaries are equitable across departments, age groups, or experience levels, and if high performers are adequately compensated. This can support compensation review and strategy development.
3. ****Data Collection Improvement****: The most immediate business insight derived from this analysis is the critical need to improve HR data collection practices, especially for `Salary` and `Performance\_Score`. Having almost half of `Salary` data missing significantly hampers any data-driven decision-making in compensation. This directly points to an operational gap that needs to be addressed.
4. ****Resource Allocation & Development****: Understanding where high performance clusters (e.g., specific departments or experience levels) can guide where to invest more resources, talent development programs, or identify areas needing intervention.


5. **Benchmarking**: With more robust data, the dataset could eventually serve as a basis for internal benchmarking of employee performance and compensation.

In conclusion, while the dataset is small and has significant missing value issues, particularly for `Salary`, it offers a starting point for understanding employee dynamics.


The first priority should be to address data quality, especially the missing values in `Salary` and `Performance\_Score`, and then proceed with the recommended preprocessing before attempting any advanced analytics or machine learning.

Visualizations


Histogram - Salary




Box Plot - Salary




Histogram - Age




Box Plot - Age




Histogram - Experience_Years




Box Plot - Experience_Years




Correlation Heatmap



Salary vs Age




Bar Chart - Department




Pie Chart - Department

Department Distribution




Bar Chart - City

City Distribution



Pie Chart - City

City Distribution



Machine Learning Recommendation

- Remove duplicate rows
- Handle missing values
- Perform feature engineering
- Encode categorical variables
- Scale numerical features if required
- Train multiple ML models
- Evaluate models using cross-validation

Generated by Autonomous Data Science Co-Pilot

Generated on: 16 July 2026 15:47

Powered by Gemini AI • Plotly • Streamlit • ReportLab